

TITLE PAGE

RetinaScan AI

Fundus photo in → HL7 FHIR R4 clinical report out · **0.59 s** · **100% local**


Problem Statement Track — [Enter track name]

Problem Statement Title — RetinaScan AI — AI-Assisted Diabetic Retinopathy Screening

PS Category — Software

Team Name — [Enter team name]

Project presentation · Siddhesh · October 2026



Real inference · Grad-CAM attention over a fundus frame

97.5% referable-DR sensitivity (95% CI 96.5–98.2)

0.59 s full pipeline on a laptop CPU · P95 0.71 s

\$0 per screening after one-time setup

One command, one laptop — a fundus photo in, a clinical report out.

PROJECT / IDEA NAME

RetinaScan AI — AI-assisted diabetic retinopathy screening & classification

ONE-LINE TAGLINE

Fundus photo in, FHIR R4 clinical report out in 0.59 s — 100% local.

PROBLEM BEING ADDRESSED

- ~77M diabetics in India; DR is the leading cause of preventable working-age blindness
- Manual grading takes 5–10 min per image at \$15–50 — specialist-bound, city-bound
- Inter-observer drift across ICDR stages; reports stay ad-hoc notes, not standards

PROPOSED SOLUTION


- Quality gate → dual-model ICDR staging + 4-lesion masks → FHIR R4 + Grad-CAM + PDF
- 4-model ConvNeXt-Tiny ensemble — 82.1% held-out; abstains below 55% confidence
- One command, fully local: 0.59 s per image on a laptop CPU — no cloud, no internet

77M

people with diabetes in India need periodic DR screening

Today: specialist-only · 5–10 min · \$15–50 · city-bound

LIVE PRODUCT — REAL MODEL OUTPUT



- ICDR stage + full probability vector
- 4 lesion masks with area %
- Abstention when confidence < 55%
- FHIR R4 DiagnosticReport + PDF
- Flagged on this run at 52.1%

TIME

5–10 min → **0.59 s**

COST

\$15–50 → **\$0**

OUTPUT

ad-hoc notes → **FHIR R4**


VERDICT

subjective → **97.5% sens**

RUNTIME

cloud → **100% local**

Quality gate → two models in parallel → one clinical report. 0.59 s end to end.




Ungradable frames are rejected with recapture guidance and low-confidence cases abstain to a clinician — the system never returns a confident wrong answer.

© Manual review recommended

Model confidence is below the abstention threshold. The stage shown below is a flagged suggestion only — a clinician must confirm before any clinical decision.

Live run: confidence 52.1% < 55% → abstention banner (real screenshot)



DEPLOYMENT TOPOLOGY

Frontend: React 19 SPA :3000

API: FastAPI :8000 · 10 endpoints

Inference: ONNX Runtime FP32 · CPU/GPU

Store: SQLite WAL · visit timeline

WHAT COMES OUT

- ICDR stage 0–4 + calibrated confidence
- 4 lesion masks with area %
- Grad-CAM attention heat-map
- HL7 FHIR R4 DiagnosticReport + PDF

STACK

Python 3.12

FastAPI

PyTorch 2.6

ONNX Runtime

OpenCV

React 19

SQLite WAL

HL7 FHIR R4

SNOMED CT

ReportLab

Both heads run on one 512² tensor: the same preprocessed frame drives ICDR staging, lesion masks, Grad-CAM explainability and the FHIR R4 report — 0.59 s on CPU, P95 0.71 s.

Evidence: 3,394 unseen images, sub-second end to end, \$0 per screen.

97.5%

referable-DR sensitivity
95% CI 96.5–98.2 · n = 3,394

0.843

ROC AUC (referable vs not)
sens 97.5% · spec 72.8%

82.1%

4-model ensemble, held-out
single best 80.4%


0.59 s

end-to-end on laptop CPU
P95 0.71 s · 100% local


EXTERNAL CONFUSION MATRIX · APTOS · N = 3,394

True stage (row = 100 %)	No DR	85% 1,363	2% 32	13% 204	0% 2	0% 4
	Mild	10% 37	0% 0	75% 269	12% 45	3% 10
	Moderate	3% 26	0% 0	54% 524	34% 333	9% 85
	Severe	1% 2	0% 0	29% 53	54% 98	15% 28
	Proliferative	3% 7	0% 1	30% 84	39% 108	28% 79
	Predicted stage					
	No DR	Mild	Moderate	Severe	Proliferative	

REFERABLE DR · 2 × 2



HELD-OUT ICDR F1 · N = 815



Per-class F1 on the 815-image held-out split. The 4-model ensemble adds +1.7 pp (80.4% → 82.1%) over the single best checkpoint.

TECHNICAL

11/11 components verified end-to-end;
ONNX FP32 on CPU

OPERATIONAL

One command: `bash start_demo.sh` ·
offline, no cloud APIs

ECONOMIC

\$0 per screen after setup vs \$15–50
manual grading

SCALABILITY

Batch triage of 20 images per call · FHIR-
ready for EHRs

Honest external read: five-class accuracy on the imbalanced APTOS set is 60.8% (95% CI 59.2–62.3) — but the decision the tool actually drives, referable vs non-referable, is right 97.5% of the time. Rare-stage recall is the next training cycle.

What changes when a screening is instant, free and offline.

1 CLINICAL CONFIDENCE

97.5% sensitivity · 97.5% NPV

2 FEWER MISSED REFERRALS

abstains < 55% → clinician review

3 NEAR-ZERO COST

\$0 per screen vs \$15–50 manual

4 MINIMAL FOOTPRINT

one laptop · offline · no cloud

5 FASTER TRIAGE

0.59 s vs 5–10 min per image

BEFORE VS AFTER · ONE SCREENING



Same laptop, same image, no internet: the marginal cost of a screen goes to zero while the wait falls from minutes to under a second.

IMPACT PATHWAY

PROBLEM

specialist-only · 5–10 min · \$15–50

SOLUTION

0.59 s local AI · FHIR R4 · abstention

IMMEDIATE IMPACT

instant triage · fewer missed referrals

LONG-TERM IMPACT

population screening · EHR-native

WHO BENEFITS

PRIMARY HEALTH CENTRES

Rural clinics screen without a specialist on site

EYE HOSPITALS

Batch triage of 20 images + Grad-CAM review queue

PUBLIC HEALTH PROGRAMMES

Population screening at near-zero marginal cost

The datasets, prior work, standards and ablations behind every claim.

RESEARCH / LITERATURE SURVEY

Porwal et al., IDRiD (Data 2018) — DR grades + pixel-level lesion masks → segmenter

APTOS 2019 (Kaggle) — 3,662 graded fundus images → train + external test

DDR dataset — lesion masks lift segmenter training 54 → 811 (×15)

Ronneberger et al., U-Net (MICCAI 2015) — encoder-decoder for the 4 lesion channels




Liu et al., ConvNeXt (CVPR 2022) — backbone upgrade: +8.6 pp over EfficientNet-B2

Grad-CAM · Focal Loss (ICCV 2017) — attention heat-maps + class-imbalance loss

COMPETITIVE LANDSCAPE

	RetinaScan AI	Cloud vision APIs	Manual grading
Runs fully offline	✓	—	✓
Abstains to a human	✓	—	✓
HL7 FHIR R4 output	✓	—	—
Grad-CAM explainable	✓	✓	✓
Cost per screen	\$0	metered	\$15–50

“—” = not documented in the vendor's public material; manual-grading cost from PROJECT_INFO.



REFERENCES

1 Ronneberger et al., U-Net, MICCAI 2015 (arXiv:1505.04597) · 2 Liu et al., A ConvNet for the 2020s, CVPR 2022 (arXiv:2201.03545) · 3 Lin et al., Focal Loss, ICCV 2017 · 4 Selvaraju et al., Grad-CAM, ICCV 2017 (arXiv:1610.02391) · 5 Porwal et al., IDRiD, Data 2018 · 6 APTOS 2019 Blindness Detection, Kaggle · 7 HL7 FHIR R4 · SNOMED CT · LOINC (hl7.org · snomed.org · loinc.org)

IMPORTANT INSTRUCTIONS

Keep all information concise and relevant.

Clearly explain the problem and proposed solution.

Use authentic research and reliable references.

Ensure technical claims are realistic and implementable.

Maintain consistency in fonts, colors, and formatting.

Use diagrams / flowcharts wherever they improve understanding.

Ensure all team / project information is accurate.

Avoid unnecessary text and decorative elements.